

# Depression Sentiment analysis with Embedding and LSTM

## Import Dependence

```
In [1]: import numpy as np
import pandas as pd
import nltk
import re
import string
from nltk.corpus import stopwords
nltk.download("stopwords")
stemmer = nltk.SnowballStemmer("english")
stopword=set(stopwords.words('english'))
from nltk.sentiment.vader import SentimentIntensityAnalyzer
from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator
```

```
[nltk_data] Downloading package stopwords to /Users/vikky/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

## Data collection

```
In [2]: #loading the dataset panda dataframe
data = pd.read_csv('depression_dataset_reddit_cleaned.csv')
```

```
In [3]: #check first five rows of the dataset
data.head()
```

```
Out[3]:
```

	clean_text	is_depression
0	we understand that most people who reply immed...	1
1	welcome to r depression s check in post a plac...	1
2	anyone else instead of sleeping more when depr...	1
3	i ve kind of stuffed around a lot in my life d...	1
4	sleep is my greatest and most comforting escap...	1

In [4]: *#check last five rows of the dataset*  
`data.tail()`

Out [4]:

	<b>clean_text</b>	<b>is_depression</b>
<b>7726</b>	is that snow	0
<b>7727</b>	moulin rouge mad me cry once again	0
<b>7728</b>	trying to shout but can t find people on the list	0
<b>7729</b>	ughh can t find my red sox hat got ta wear thi...	0
<b>7730</b>	slept wonderfully finally tried swatching for ...	0

In [5]: *#check columns of the dataset*  
`data.columns`

Out [5]: Index(['clean\_text', 'is\_depression'], dtype='object')

In [6]: *#check more infomation of the dataset*  
`data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7731 entries, 0 to 7730
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   clean_text      7731 non-null   object
1   is_depression   7731 non-null   int64
dtypes: int64(1), object(1)
memory usage: 120.9+ KB
```

In [7]: *#check mathmatic realtionship of the dataset*  
`data.describe()`

Out [7]:

	<b>is_depression</b>
<b>count</b>	7731.000000
<b>mean</b>	0.495537
<b>std</b>	0.500012
<b>min</b>	0.000000
<b>25%</b>	0.000000
<b>50%</b>	0.000000
<b>75%</b>	1.000000
<b>max</b>	1.000000

```
In [8]: #check missing value of the dataset
data.isnull().sum()
```

```
Out[8]: clean_text      0
is_depression      0
dtype: int64
```

## EDA

```
In [9]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [10]: data['is_depression'].value_counts()
```

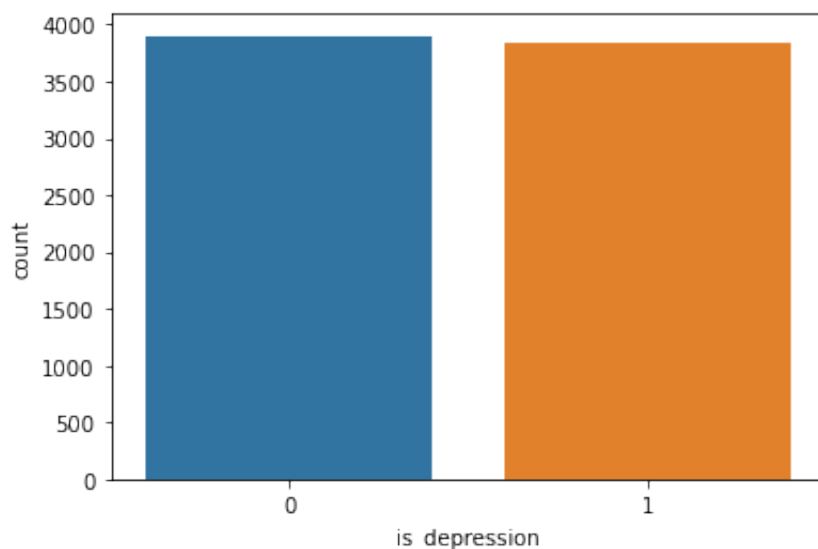
```
Out[10]: 0    3900
1    3831
Name: is_depression, dtype: int64
```

```
In [11]: #count the value is_depression dataset columns
sns.countplot(data['is_depression'])
```


/Users/vikky/opt/anaconda3/lib/python3.9/site-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[11]: <AxesSubplot:xlabel='is_depression', ylabel='count'>
```



```
In [15]: #Now let's have a look at the kind of words people use in the is_de
text = " ".join(i for i in data.clean_text)
stopwords = set(STOPWORDS)
wordcloud = WordCloud(stopwords=stopwords, background_color="white")
plt.figure(figsize=(15,10))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.show()
```



## Page 4 of 12

```
In [16]: #splitting the dataset in x and y
x = data["clean_text"]
y = data["is_depression"]
```

```
In [17]: #print x and y
print(x)
print(y)
```

```
0      understand peopl repli immedi op invit talk pr...
1      welcom r depress check post place take moment ...
2      anyon els instead sleep depress stay night avo...
3      kind stuf around lot life delay inevit work jo...
4      sleep greatest comfort escap whenev wake day l...
...
7726                                     snow
7727                                moulin roug mad cri
7728                                tri shout find peopl list
7729      ughh find red sox hat got ta wear creepi nick ...
7730      slept wonder final tri swatch new project clas...
Name: clean_text, Length: 7731, dtype: object
0      1
1      1
2      1
3      1
4      1
...
7726      0
7727      0
7728      0
7729      0
7730      0
Name: is_depression, Length: 7731, dtype: int64
```

## Train-Test split

```
In [18]: #splitting the dataset in X_train and Y_train
from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(x,y,test_size=0.2)
```

```
In [19]: #check shape X_train and Y_train
print(X_train.shape,X_test.shape,Y_train.shape,Y_test.shape)

(6184,) (1547,) (6184,) (1547,)
```

## Text Vectorization

```
In [20]: #using Text Vectorization
round(sum([len(i.split()) for i in X_train])/len(X_train))
```

Out[20]: 34

```
In [21]: # import TextVectorization
import tensorflow as tf
from tensorflow.keras.layers import TextVectorization
```

```
In [22]: # splitting the dataset in text_vectorizer
max_vocab_length = 10000
max_length = 34

text_vectorizer = TextVectorization(max_tokens=max_vocab_length,
                                     output_mode="int",
                                     output_sequence_length=max_length)
```

```
In [23]: #check X_train after splitting Text_vectorizer
text_vectorizer.adapt(X_train)
```

```
In [24]: #find Words in vocab
words_in_vocab = text_vectorizer.get_vocabulary()
top_5_words = words_in_vocab[:5]
bottom_5_words = words_in_vocab[-5:]
print(f"Vocabulary size: {len(words_in_vocab)}")
print(f"Top 5 most common words: {top_5_words}")
print(f"Bottom 5 least common words: {bottom_5_words}")
```

Vocabulary size: 10000

Top 5 most common words: ['', '[UNK]', 'feel', 'wa', 'like']

Bottom 5 least common words: ['inna', 'inlov', 'inlaw', 'inkart', 'ink']

## Embedding Layer

```
In [25]: #using Embedding layer
from tensorflow.keras import layers

embedding = layers.Embedding(input_dim=max_vocab_length, output_dim=
                             embeddings_initializer="uniform",
                             input_length=max_length
                             )
```

```
In [26]: from tensorflow.keras import layers
inputs = layers.Input(shape=(1,), dtype=tf.string)
x = text_vectorizer(inputs)
x = embedding(x)
x = layers.GlobalAveragePooling1D()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model_1 = tf.keras.Model(inputs, outputs, name="model_1_dense")
```

```
In [27]: #model compile
model_1.compile(loss='binary_crossentropy', optimizer='adam', metrics
```

```
In [28]: #train model
history = model_1.fit(X_train, Y_train, validation_data=(X_test, Y_test))
```

```
Epoch 1/10
194/194 [=====] - 8s 27ms/step - loss: 0.4695 - accuracy: 0.8357 - val_loss: 0.3507 - val_accuracy: 0.8655
Epoch 2/10
194/194 [=====] - 5s 25ms/step - loss: 0.2837 - accuracy: 0.8991 - val_loss: 0.2763 - val_accuracy: 0.8966
Epoch 3/10
194/194 [=====] - 5s 25ms/step - loss: 0.2158 - accuracy: 0.9221 - val_loss: 0.2265 - val_accuracy: 0.9147
Epoch 4/10
194/194 [=====] - 5s 27ms/step - loss: 0.1663 - accuracy: 0.9402 - val_loss: 0.1883 - val_accuracy: 0.9276
Epoch 5/10
194/194 [=====] - 6s 31ms/step - loss: 0.1296 - accuracy: 0.9560 - val_loss: 0.1622 - val_accuracy: 0.9412
Epoch 6/10
194/194 [=====] - 5s 27ms/step - loss: 0.1046 - accuracy: 0.9680 - val_loss: 0.1458 - val_accuracy: 0.9502
Epoch 7/10
194/194 [=====] - 6s 30ms/step - loss: 0.0867 - accuracy: 0.9753 - val_loss: 0.1340 - val_accuracy: 0.9606
Epoch 8/10
194/194 [=====] - 5s 28ms/step - loss: 0.0738 - accuracy: 0.9788 - val_loss: 0.1314 - val_accuracy: 0.9580
Epoch 9/10
194/194 [=====] - 6s 28ms/step - loss: 0.0638 - accuracy: 0.9821 - val_loss: 0.1271 - val_accuracy: 0.9599
Epoch 10/10
194/194 [=====] - 6s 29ms/step - loss: 0.0569 - accuracy: 0.9843 - val_loss: 0.1283 - val_accuracy: 0.9580
```

In [30]: `history.history`

```
Out[30]: {'loss': [0.46946293115615845,
0.2837080955505371,
0.21580171585083008,
0.16626225411891937,
0.12964126467704773,
0.10459963977336884,
0.08674520999193192,
0.07382360845804214,
0.06380867213010788,
0.05686891824007034],
'accuracy': [0.835705041885376,
0.8990944623947144,
0.9220569133758545,
0.9401682019233704,
0.9560155272483826,
0.9679818749427795,
0.9752587080001831,
0.9788162708282471,
0.9820504784584045,
0.9843143820762634],
'val_loss': [0.3507160246372223,
0.2762649357318878,
0.2264728993177414,
0.18832530081272125,
0.16224391758441925,
0.14581571519374847,
0.13403543829917908,
0.13141471147537231,
0.12711641192436218,
0.12834328413009644],
'val_accuracy': [0.8655462265014648,
0.8965740203857422,
0.9146735668182373,
0.9276018142700195,
0.9411764740943909,
0.9502262473106384,
0.9605688452720642,
0.9579831957817078,
0.9599224328994751,
0.9579831957817078]}
```

## Classification Report for Model 1

In [32]: `Y_pred=model_1.predict(X_test)`  
`Y_pred=(Y_pred>=0.5).astype("int")`

49/49 [=====] - 0s 2ms/step



```
In [33]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
print(classification_report(Y_test,Y_pred))
```

	precision	recall	f1-score	support
0	0.93	0.98	0.96	770
1	0.98	0.93	0.96	777
accuracy			0.96	1547
macro avg	0.96	0.96	0.96	1547
weighted avg	0.96	0.96	0.96	1547

## LSTM model

```
In [34]: #using Lstm model
from tensorflow.keras import layers
inputs = layers.Input(shape=(1,), dtype=tf.string)
x = text_vectorizer(inputs)
x = embedding(x)
x = layers.LSTM(64, activation="tanh")(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model_2 = tf.keras.Model(inputs, outputs, name="model_2_lstm")
```

```
In [35]: #compile model
model_2.compile(loss='binary_crossentropy',optimizer='adam',metrics
```

```
In [36]: #train model using lstm
history = model_2.fit(X_train,Y_train,validation_data=(X_test,Y_test))

Epoch 1/10
194/194 [=====] - 17s 67ms/step - loss: 0.1219 - accuracy: 0.9657 - val_loss: 0.1875 - val_accuracy: 0.9418
Epoch 2/10
194/194 [=====] - 13s 67ms/step - loss: 0.0443 - accuracy: 0.9880 - val_loss: 0.1772 - val_accuracy: 0.9463
Epoch 3/10
194/194 [=====] - 13s 67ms/step - loss: 0.0337 - accuracy: 0.9898 - val_loss: 0.1912 - val_accuracy: 0.9535
Epoch 4/10
194/194 [=====] - 12s 61ms/step - loss: 0.0386 - accuracy: 0.9905 - val_loss: 0.1657 - val_accuracy: 0.9489
Epoch 5/10
194/194 [=====] - 12s 61ms/step - loss: 0.0285 - accuracy: 0.9927 - val_loss: 0.1657 - val_accuracy: 0.9457
Epoch 6/10
194/194 [=====] - 12s 60ms/step - loss: 0.0175 - accuracy: 0.9958 - val_loss: 0.2598 - val_accuracy: 0.9438
Epoch 7/10
194/194 [=====] - 12s 61ms/step - loss: 0.0176 - accuracy: 0.9947 - val_loss: 0.2520 - val_accuracy: 0.9438
Epoch 8/10
194/194 [=====] - 12s 61ms/step - loss: 0.0223 - accuracy: 0.9930 - val_loss: 0.2109 - val_accuracy: 0.9483
Epoch 9/10
194/194 [=====] - 15s 78ms/step - loss: 0.0133 - accuracy: 0.9963 - val_loss: 0.3077 - val_accuracy: 0.9302
Epoch 10/10
194/194 [=====] - 13s 68ms/step - loss: 0.0362 - accuracy: 0.9897 - val_loss: 0.2668 - val_accuracy: 0.9392
```

In [37]: `history.history`

```
Out[37]: {'loss': [0.12194930762052536,
0.04434535279870033,
0.03370266035199165,
0.03857898712158203,
0.028453631326556206,
0.01745663397014141,
0.01756499893963337,
0.022268325090408325,
0.013344790786504745,
0.0361902117729187],
'accuracy': [0.9657179713249207,
0.988033652305603,
0.989812433719635,
0.9904592633247375,
0.9927231669425964,
0.9957956075668335,
0.994663655757904,
0.9930465817451477,
0.9962807297706604,
0.9896507263183594],
'val_loss': [0.18753276765346527,
0.17715151607990265,
0.19118280708789825,
0.16568949818611145,
0.1657085120677948,
0.2597535252571106,
0.2520136535167694,
0.21087147295475006,
0.3077019453048706,
0.26683393120765686],
'val_accuracy': [0.94182288646698,
0.9463477730751038,
0.953458309173584,
0.9489334225654602,
0.9457013607025146,
0.9437621235847473,
0.9437621235847473,
0.9482870101928711,
0.930187463760376,
0.9392372369766235]}
```

## Classification Report for Model 2

In [38]: `Y_pred=model_2.predict(X_test)`  
`Y_pred=(Y_pred>=0.5).astype("int")`

49/49 [=====] - 1s 11ms/step

```
In [39]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
print(classification_report(Y_test,Y_pred))
```

	precision	recall	f1-score	support
0	0.95	0.93	0.94	770
1	0.93	0.95	0.94	777
accuracy			0.94	1547
macro avg	0.94	0.94	0.94	1547
weighted avg	0.94	0.94	0.94	1547

```
In [ ]:
```